


Qwen2.5-1.5B-Instruct · NVIDIA GeForce RTX 4090
vLLM 0.10.2 · 128-token fixed decode · warm synthetic workload

--- Prefix cache OFF — Prefix cache ON ● Replicate 1 ■ Replicate 2 ▲ Replicate 3

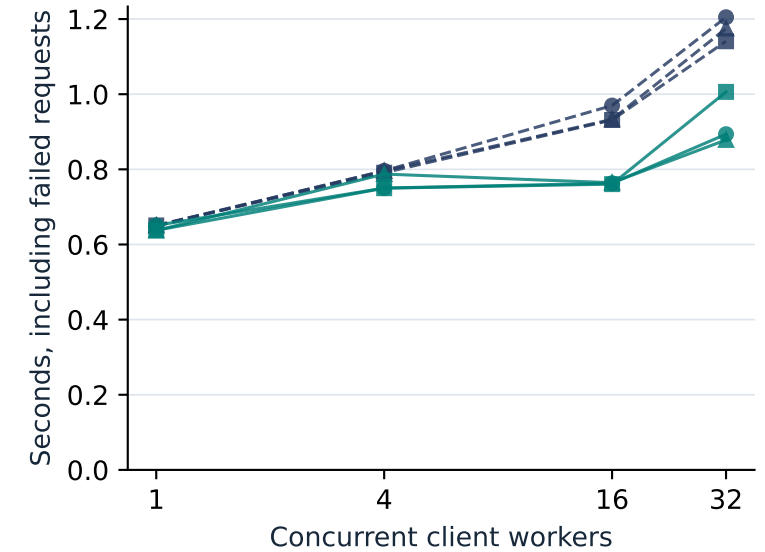
Decode throughput



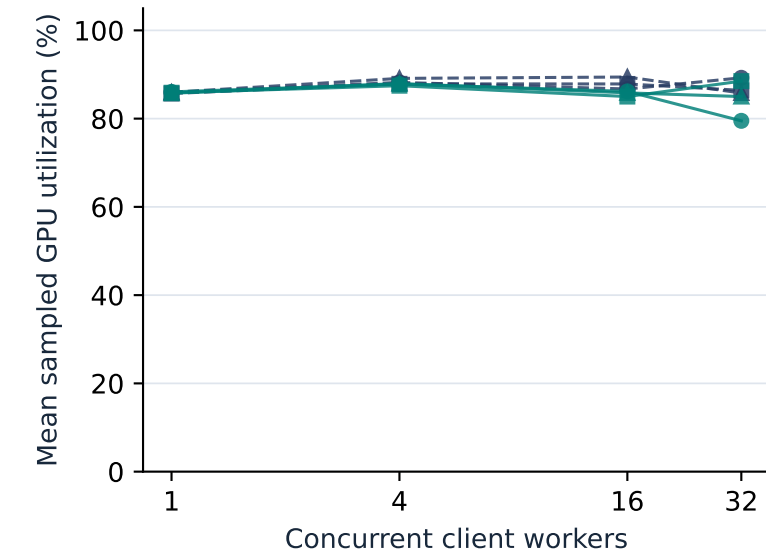
Time to first text · p95



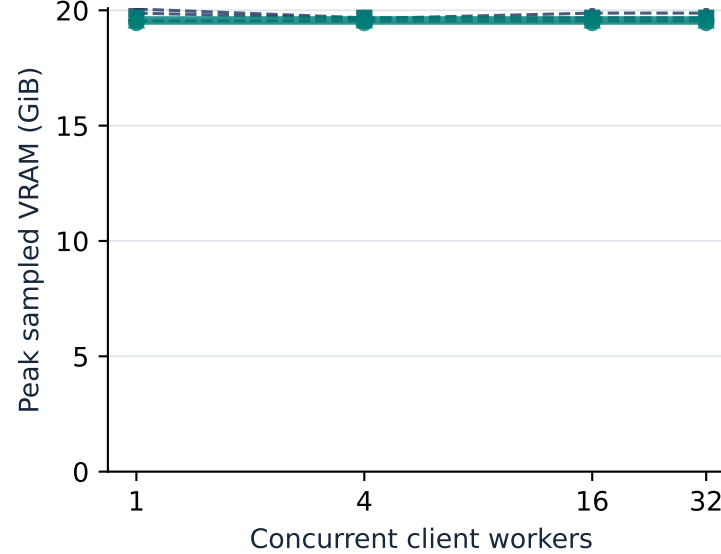
Request latency · p95



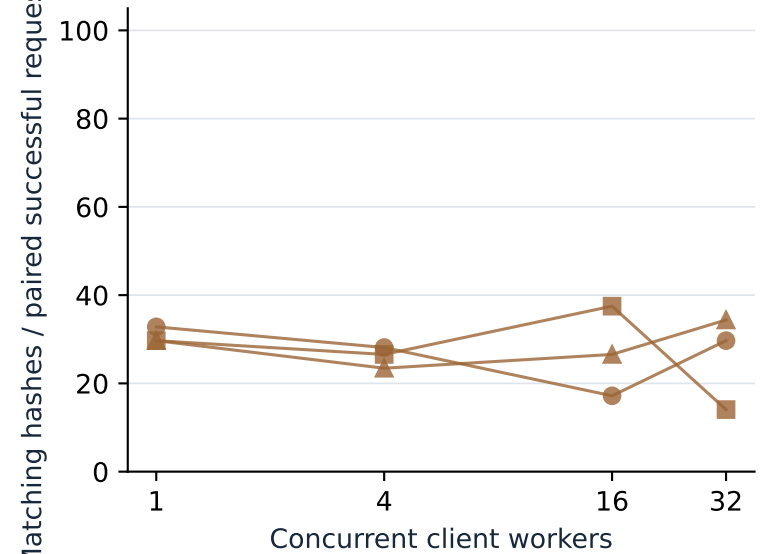
Device utilization



Device memory



Cache modes · exact text agreement



1,536 measured requests; 0 failures; 0 output-length mismatches. Every replicate shown; no population confidence intervals.
p95 values are sample percentiles. Text agreement is not answer quality. VRAM includes allocated KV cache.